
Trajectory Supervision for Continual Tool-Use Learning in LLMs

Vishnu Vardhan Reddy B
CS 590NN
UMass Amherst
vbheemreddy@umass.edu

Sagnik Chatterjee
CS 590NN
UMass Amherst
sagnikchatte@umass.edu

Soumik Bhatta
CS 590NN
UMass Amherst
sbhatta@umass.edu

Abstract

Most language-model training data shows final artifacts, not the process that produced them. We study a tractable version of this question in tool use: when a model learns a stream of new API domains, does keeping tool-use trajectories help compared with stripping the intermediate API trace? We fine-tune Llama 3.1 8B Instruct with QLoRA on API-Bank using four sequential domain blocks. Condition A strips previous API request/response lines from the prompt and trains the model to predict the next API call. Condition B keeps the trajectory context. In a single-seed pilot, full held-out generation evaluation shows that Condition B reaches 56.9% final exact full-call accuracy compared with 39.2% for Condition A. B also improves final API-name accuracy by 7.7 points. The result is promising but not conclusive: B uses 25.1% more training tokens, the run uses one seed, and the task is next-call prediction rather than full dialogue success.

1 Introduction

Language models are usually trained on finished text: answers, articles, code, and other final products. The intermediate process that produced the answer is usually missing. That gap motivates process-supervision work in reasoning (Lightman et al., 2023), but it is hard to test directly because real human thought traces are private, expensive, and often unavailable.

Tool-use data gives a narrower but concrete proxy. A tool-use example can include a user request, an API call, the API response, and the next assistant action. These traces are not human internal reasoning. They are external interaction records. Still, they expose an action-observation sequence that final-response data removes. Our project asks whether this extra trajectory signal changes how a model adapts when tool domains arrive sequentially.

We use continual learning as a stress test, not as a new algorithmic contribution. If two supervision formats perform similarly on one domain, differences may become clearer when the model learns a sequence of new domains and must retain older ones. We therefore split API-Bank (Li et al., 2023) into four sequential blocks and evaluate after each training block on all blocks.

The novel part of this project is the controlled pilot framing: we compare a stripped-context next-API-call baseline against a trajectory-context baseline under the same model, seed, stream ordering, QLoRA setup, and evaluation code. We do not claim to have solved catastrophic forgetting or to have shown that API traces work like human thought. The contribution is narrower: a first controlled measurement of whether retaining tool-use trajectory information changes continual tool-call learning.

Our main findings are:

- Full held-out generation evaluation favors trajectory context. After training on all four blocks, Condition B reaches 56.9% exact full-call accuracy, compared with 39.2% for Condition A.

- The largest improvement is in selecting the right tool. B reduces final wrong-API errors from 102 to 12, while exact parameter generation remains imperfect.
- The evidence is limited by one seed and a token-budget confound. B trains on 2.32M tokens versus 1.86M for A, so the result supports a hypothesis but does not isolate the cause.

2 Related work

Process supervision. Lightman et al. compare process and outcome supervision for mathematical reasoning and show that step-level supervision can be more effective than final-answer supervision (Lightman et al., 2023). Our work is inspired by that distinction, but our data is different. API trajectories are not annotated reasoning steps; they are records of tool calls and tool observations. This makes the supervision weaker as a cognitive model but easier to collect in real systems.

Tool-use learning. Toolformer showed that language models can learn to call external tools from self-supervised tool-use traces (Schick et al., 2023). API-Bank provides a compact benchmark for tool-augmented language models with API calls, responses, and dialogue context (Li et al., 2023). ToolLLM and ToolBench scale this idea to many real-world APIs and solution paths (Qin et al., 2023). We use API-Bank because it is small enough for repeated QLoRA experiments in a course project while still containing structured tool-call supervision.

Continual learning in LLMs. Sequential fine-tuning can damage previously learned behavior. TRACE formalizes continual learning evaluation for LLMs and reports severe forgetting under sequential training (Wang et al., 2023). We borrow the evaluation lens: after training on block D_i , evaluate on all blocks D_j . Our aim is not to introduce a new forgetting method such as replay or regularization. Instead, we ask whether the supervision format itself changes adaptation and retention.

Efficient adaptation. LoRA trains low-rank adapter weights instead of updating all parameters (Hu et al., 2022). QLoRA adds 4-bit quantization to make fine-tuning larger models feasible on limited hardware (Dettmers et al., 2023). We use QLoRA with Llama 3.1 8B Instruct (Dubey et al., 2024) so that both conditions can be trained and checkpointed in Colab.

3 Method

3.1 Task and data

We use API-Bank and construct a stream of four domain blocks, D_1 through D_4 . Each block contains a disjoint set of API examples. Training proceeds sequentially: train on D_1 , evaluate on all blocks, then train on D_2 , and so on. This produces a 4×4 evaluation matrix. Rows are training stage and columns are evaluation block.

The full held-out generation evaluation scores only examples whose expected output contains a parseable API call. The final-stage scored evaluation totals are 126, 104, 103, and 107 examples for blocks D_1 through D_4 , respectively. The training notebooks also run a faster sampled evaluation with 32 examples per block after each stage.

3.2 Conditions

Condition A is the stripped-context baseline. It removes prior API-Request and API-Response lines from the input context before training and generation evaluation. It still predicts the next API call when the target output is an API call. This is narrower than the original proposal’s broad “outcome-only” framing, so we describe it as stripped-context next-call supervision.

Condition B is the trajectory-context condition. It keeps the previous API request and response lines in the prompt. This gives the model access to the action-observation sequence before the next API call. Both conditions use the same base model, stream ordering, seed, optimizer settings, and scoring code.

Table 1: Sampled generation evaluation from the training notebook. Scores are exact full-call accuracy unless noted.

Condition	Final AA	BWT	FWT	Avg. forgetting	AULC	Train tokens
A: stripped context	38.3	-10.4	22.9	10.4	41.8	1.86M
B: trajectory context	53.9	-13.5	33.3	13.5	57.2	2.32M

3.3 Model and training

All runs use meta-llama/Llama-3.1-8B-Instruct. We fine-tune with QLoRA using 4-bit NF4 quantization, LoRA rank 32, LoRA alpha 64, bfloat16 computation, maximum sequence length 1024, and an effective batch size of 16. We use AdamW with a learning rate of 2×10^{-4} . Each block is trained for three epochs. The training runs use seed 42.

The main fairness limitation is token count. Condition A consumes 1,857,169 training tokens; Condition B consumes 2,324,314 training tokens. B therefore sees 25.1% more tokens because trajectory context is longer. We report this explicitly and treat token-matched training as future work.

3.4 Evaluation metrics

For generation, we greedily decode the model’s next action and parse API calls with a regular expression of the form `[ApiName(param='value')]`. We report:

- **API-name accuracy:** the generated API name matches the expected API name.
- **Exact full-call accuracy:** the generated API name and normalized parameter dictionary exactly match the expected call.
- **Name-plus-any-param accuracy:** the API name is correct and at least one expected parameter-value pair is correct.
- **Malformed/no-call rate:** the generated text does not contain a parseable API call.

For the sampled continual-learning analysis, we also compute average accuracy (AA), backward transfer (BWT), forward transfer (FWT), average forgetting, and area under the learning curve (AULC) using exact full-call accuracy.

4 Results

4.1 Sampled continual-learning evaluation

Table 1 summarizes the 32-example-per-block evaluation produced during training. B has higher final average accuracy and higher forward transfer, but it also has more negative BWT and larger average forgetting under this sampled metric. This means B is not simply better on every continual-learning statistic; it learns the next-call format more effectively, but it can still lose performance on earlier blocks.

4.2 Full held-out generation evaluation

After training, we ran a separate full generation evaluation over every scored held-out example for each saved adapter. This removes the 32-example sampling limit used during training. Figure 2 shows the exact full-call accuracy matrices.

The final-stage comparison is the clearest result. Condition B reaches 56.9% mean exact full-call accuracy, compared with 39.2% for Condition A, a difference of 17.7 percentage points. On API-name accuracy, B reaches 74.3% compared with 66.6% for A, a difference of 7.7 points. Table 2 shows the final-stage block-level results.




Figure 1: Sampled evaluation used during training. Condition B obtains lower loss and higher full-call accuracy in the sampled analysis, while both conditions show nonzero forgetting.




Figure 2: Full held-out exact full-call accuracy. Rows are training stages and columns are evaluation blocks. Condition B is higher on every final-stage block and most earlier stages.

4.3 Error analysis

The error categories clarify what trajectory context changed. At the final stage, A produced 102 wrong-API errors, while B produced only 12. This supports the interpretation that trajectory context helps the model choose the correct tool. However, B also produced 101 malformed/no-call errors, compared with 45 for A. The B model often had the right tool structure at the name level but did not always produce an exactly parseable or parameter-complete call.

Table 2: Full held-out final-stage generation results after training through D_4 .

Condition	Metric	D1	D2	D3	D4	Mean
A	Exact full-call	35.7	43.3	32.0	45.8	39.2
B	Exact full-call	57.9	61.5	44.7	63.6	56.9
A	API-name	64.3	62.5	60.2	79.4	66.6
B	API-name	73.8	82.7	67.0	73.8	74.3
A	Name + any param	51.6	56.7	50.5	65.4	56.1
B	Name + any param	67.5	76.9	61.2	72.0	69.4




Figure 3: Final-stage full evaluation by block. B improves exact full-call and name-plus-any-param accuracy on every block.

5 Discussion

The main result supports the project hypothesis in a limited form: keeping trajectory context improves final held-out next-API-call generation in this single-seed API-Bank stream. The improvement is strongest for tool selection. The action-observation history likely gives the model useful structure for choosing which API to call next.

That said, this result does not prove that process supervision is inherently better. The most obvious confound is token budget: B sees 25% more tokens, and more data can help even if the extra content is not causally important. We also have only one seed, so we cannot report confidence intervals. API-Bank traces are structured and partly synthetic, which makes them a weaker stand-in for real reasoning data. And our task is next API-call prediction, not full end-to-end task success—a narrower evaluation than what a deployment setting would require.

The malformed/no-call result is also important. B makes fewer wrong tool choices, but it makes more parse failures. Longer trajectory prompts may help the model identify the right tool while making exact output formatting harder to get right. Our parser is also strict: a call with the correct API and nearly correct parameters still counts as a failure if the string does not exactly match. A semantic parameter scorer alongside exact-match scoring would help separate these two effects.

Table 3: Final-stage full-eval error category counts across all scored blocks.

Category	A	B
Exact full call	172	251
Correct API, some params	74	54
Correct API, wrong params	47	22
Wrong API	102	12
Malformed or no call	45	101

6 Conclusion and future work

We tested whether trajectory context changes continual tool-use learning in a controlled Llama 3.1 8B QLoRA pilot. The full held-out evaluation shows a clear advantage for trajectory context on final exact API-call generation: 56.9% versus 39.2%. The result is promising, but it remains a pilot because the trajectory condition uses more tokens and the experiment uses one seed.

The most urgent follow-up is running multiple seeds and reporting confidence intervals, since a single-seed result cannot support strong claims. A token-matched variant of Condition B would also help separate the effect of trajectory content from the effect of seeing more tokens. Beyond that, longer task streams and a semantic parameter scorer—one that distinguishes formatting errors from genuinely wrong tool calls—would strengthen the evaluation. Replay or retrieval methods are worth exploring, though they would be new continual-learning interventions rather than part of the A/B supervision question.

7 Contributions

All three students contributed to the project. Individual responsibilities are listed in Table 4; shared work included project framing, result interpretation, and report writing.

Table 4: Equal project contribution summary.

Student	Contributions
Vishnu Vardhan Reddy B	Data preprocessing pipeline, evaluation and scoring code, artifact validation, full generation evaluation runs. Shared: project framing, result interpretation, report writing.
Sagnik Chatterjee	Training notebook implementation, continual-learning metric computation, Colab experiment execution. Shared: project framing, result interpretation, report writing.
Soumik Bhatta	Error analysis, figure generation, presentation preparation, experiment verification. Shared: project framing, result interpretation, report writing.

References

- Tim Dettmers, Artidoro Pagnoni, Ari Holtzman, and Luke Zettlemoyer. QLoRA: Efficient finetuning of quantized LLMs. *NeurIPS*, 2023.
- Abhimanyu Dubey et al. The Llama 3 herd of models. *arXiv:2407.21783*, 2024.
- Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. LoRA: Low-rank adaptation of large language models. *ICLR*, 2022.
- Minghao Li, Yingxiu Zhao, Bowen Yu, Feifan Song, Hangyu Li, Haiyang Yu, Zhoujun Li, Fei Huang, and Yongbin Li. API-Bank: A comprehensive benchmark for tool-augmented LLMs. *arXiv:2304.08244*, 2023.
- Hunter Lightman et al. Let’s verify step by step. *arXiv:2305.20050*, 2023.
- Yujia Qin et al. ToolLLM: Facilitating large language models to master 16000+ real-world APIs. *arXiv:2307.16789*, 2023.

Timo Schick, Jane Dwivedi-Yu, Roberto Dessi, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. *NeurIPS*, 2023.

Xiao Wang et al. TRACE: A comprehensive benchmark for continual learning in large language models. *arXiv:2310.06762*, 2023.